

1 Horn-SAT and Renamable Horn-SAT

For parts (a) and (b), let us first establish that Horn clauses can be rewritten as implications. We show this by case analysis on the structure of Horn clauses.

A Horn clause with exactly one positive literal

In the case of a Horn clause with exactly one positive literal, e.g. $(\neg x_1 \vee \neg x_2 \vee \dots \vee x_k)$, material implication (the rule of inference) allows us to rewrite this clause to $(x_1 \wedge x_2 \wedge \dots \Rightarrow x_k)$. The single positive literal clause (x) is a special case, whose implication can be thought of as $(true \Rightarrow x)$.

A Horn clause with no positive literals

In the case of a Horn clause with no positive literals, e.g. $(\neg x_1 \vee \neg x_2 \vee \dots \vee \neg x_k)$, we can think of its implication as $(x_1 \wedge x_2 \wedge \dots \wedge x_k \Rightarrow false)$. This holds for the single negative literal clause (x) , i.e. $(\neg x \Rightarrow false)$.

1.1 (a) A Linear Time Algorithm for Deciding the Satisfiability of HornSAT Formulae

Assumptions and Data Structures

Let us assume that we have a Horn-SAT formula, ϕ , in CNF. To begin, we initialize a collection of data structures to assist with tracking assignments to literals.

1. For each positive literal p in ϕ , we compute a list cl containing those clauses in which p appears as a negative literal.
2. We also compute an array nls such that $nls[c]$ returns the number of negative literals in clause c that have current truth value *false* (0). A clause c can be processed for assignment by our algorithm if $nls[c] = 0$; in other words, its positive literal *must* be *true* (1).
3. We also create an array pls such that $pls[c]$ returns the positive literal in clause c , if one exists.
4. Last, we instantiate a queue q that will contain clauses that are ready to be processed (those for which $nls[c] = 0$ as stated above).

Procedure

To begin, we initialize q to contain those clauses that are comprised of a single positive literal. In implication terms, we can think of these clauses as those of the form $(true \Rightarrow x)$. We *know* that these positive literals must eventually be assigned the value *true* (1) if ϕ is satisfiable; moreover, $nls[c]$ of these clauses, by definition, evaluates to 0. We then enter a **while** loop guarded by the condition that q is not empty, that is, that there are still some clauses for which $nls[c] \neq 0$.

Within the body of the **while** loop, we enter a **for** loop iterating from 0 to the length of the queue, which is initially equivalent to the number of positive unit clauses in ϕ . In the body of the **for** loop, we pop off the head of the queue and store its value in a variable c_1 . We then access the positive literal p_1 in c_1 by indexing into pls ($pls[c_1]$) and set its value to *true*. We also store p_1 in a variable named *next*.

We then use *next* to index into ϕ and obtain the clause list cl related to c_1 . Concretely, cl represents the list of all clauses c' in which p_1 (now set to *true* (1)) appears as a negative literal. We “remove” this literal from these clauses by setting $nls[c'] = nls[c'] - 1$, in essence decrementing the count of negative literals in c' that have current truth assignment *false* (0).

We then check in a conditional if $nls[c'] = 0$, meaning that c' is now ready to be processed. If so, we use c' to index into pls ($pls[c']$) to find its positive literal p' . If the positive literal is not yet assigned a truth value, then we set its current assignment to *true* (1) and add c' to q to be processed in a later iteration of the **while** loop. If the positive literal is assigned *false* (0), then we've hit a clause for which all negative literals have current truth value *true*(1) and its positive literal is *false* (0). Therefore, we return UNSAT. The **while** loop continues to process entries in the queue until it is empty and, if this point is reached, we return SAT.

Algorithm 1 Linear Time HornSAT Satisfiability

```

1: procedure LINEARHORN( $\phi$ )
2:   for all  $p$  in  $\phi$  do
3:     Initialize:
4:      $p.cl \leftarrow list[c \text{ s.t. } p \text{ is a negative literal in } c]$ 
5:   end for
6:   Initialize:
7:    $nls \leftarrow array[0..number \text{ of clauses in } \phi] \text{ of } 0..max \text{ num of literals}$ 
8:    $pls \leftarrow array[0..number \text{ of clauses in } \phi] \text{ of } 0..max \text{ num of literals}$ 
9:    $q \leftarrow \text{queue of clauses comprised of a single positive literal}$ 
10:
11:  while  $q.length \neq 0$  do
12:    for  $i = 0$  to  $q.length$  do
13:       $c_1 \leftarrow q.pop()$ 
14:       $pls[c_1] \leftarrow true; next \leftarrow pls[c_1]$ 
15:
16:      for all  $c'$  in  $\phi[next].cl$  do ▷ Iterate over all clauses in which  $p_1$  is a negative literal
17:         $nls[c'] \leftarrow nls[c'] - 1$ 
18:        if  $nls[c'] = 0$  then
19:           $p' \leftarrow pls[c']$ 
20:          if  $p'$  exists then
21:            if  $p' \neq false$  then
22:              Set  $p'$  to true
23:               $q.push(c')$ 
24:            else
25:              return UNSAT
26:            end if
27:          end if
28:        end if
29:      end for
30:    end for
31:  end while
32:
33:  return SAT
34: end procedure

```

1.2 (b) A Polynomial Time Algorithm for Deciding Whether a CNF Formula is Renamable Horn

We start by giving a description of the logical procedure for determining whether or not a formula, F , is renamable Horn before giving the precise algorithm.

Let F be represented as a set of clauses $F = \{C_1, \dots, C_m\}$, with each clause C_i composed of a set of literals $\{L_{i1}, \dots, L_{il}\}$. A *renaming* of F would entail replacing all literals in F whose variable n also appears in a separate set $A = \{n_1, \dots, n_n\}$